



Generic XDP

From: David Miller <davem-AT-davemloft.net>
To: netdev-AT-vger.kernel.org
Subject: [PATCH v3 net-next RFC] Generic XDP
Date: Wed, 12 Apr 2017 14:54:15 -0400 (EDT)
Message-ID: <20170412.145415.1441440342830198148.davem@davemloft.net>
Cc: xdp-newbies-AT-vger.kernel.org
Archive-link: [Article](#)

This provides a generic SKB based non-optimized XDP path which is used if either the driver lacks a specific XDP implementation, or the user requests it via a new IFLA_XDP_FLAGS value named XDP_FLAGS_SKB_MODE.

It is arguable that perhaps I should have required something like this as part of the initial XDP feature merge.

I believe this is critical for two reasons:

- 1) Accessibility. More people can play with XDP with less dependencies. Yes I know we have XDP support in virtio_net, but that just creates another dependency for learning how to use this facility.

I wrote this to make life easier for the XDP newbies.

- 2) As a model for what the expected semantics are. If there is a pure generic core implementation, it serves as a semantic example for driver folks adding XDP support.

This is just a rough draft and is untested.

One thing I have not tried to address here is the issue of XDP_PACKET_HEADROOM, thanks to Daniel for spotting that. It seems incredibly expensive to do a `skb_cow(skb, XDP_PACKET_HEADROOM)` or whatever even if the XDP program doesn't try to push headers at all. I think we really need the verifier to somehow propagate whether certain XDP helpers are used or not.

Signed-off-by: David S. Miller <davem@davemloft.net>

--

v3:

- Make sure XDP program sees packet at MAC header, push back MAC header if we do XDP_TX. (Alexei)
- Elide GRO when generic XDP is in use. (Alexei)
- Add XDP_FLAG_SKB_MODE flag which the user can use to request generic XDP even if the driver has an XDP implementation. (Alexei)
- Report whether SKB mode is in use in `rtnl_xdp_fill()` via XDP_FLAGS attribute. (Daniel)

v2:

- Add some "fall through" comments in switch statements based upon feedback from Andrew Lunn
- Use RCU for generic xdp_prog, thanks to Johannes Berg.

diff --git a/include/linux/netdevice.h b/include/linux/netdevice.h
index b0aa089..071a58b 100644

```

--- a/include/linux/netdevice.h
+++ b/include/linux/netdevice.h
@@ -1891,9 +1891,17 @@ struct net_device {
    struct lock_class_key    *qdisc_tx_busylock;
    struct lock_class_key    *qdisc_running_key;
    bool                      proto_down;
+   struct bpf_prog __rcu    *xdp_prog;
};
#define to_net_dev(d) container_of(d, struct net_device, dev)

+static inline bool netif_elide_gro(const struct net_device *dev)
+{
+   if (!(dev->features & NETIF_F_GRO) || dev->xdp_prog)
+       return true;
+   return false;
+}
+
#define NETDEV_ALIGN          32

static inline
diff --git a/include/uapi/linux/if_link.h b/include/uapi/linux/if_link.h
index 8b405af..633aa02 100644
--- a/include/uapi/linux/if_link.h
+++ b/include/uapi/linux/if_link.h
@@ -887,7 +887,9 @@ enum {
    /* XDP section */

#define XDP_FLAGS_UPDATE_IF_NOEXIST    (1U << 0)
-#define XDP_FLAGS_MASK                (XDP_FLAGS_UPDATE_IF_NOEXIST)
+#define XDP_FLAGS_SKB_MODE            (2U << 0)
+#define XDP_FLAGS_MASK                (XDP_FLAGS_UPDATE_IF_NOEXIST | \
+   XDP_FLAGS_SKB_MODE)

enum {
    IFLA_XDP_UNSPEC,
diff --git a/net/core/dev.c b/net/core/dev.c
index ef9fe60e..9ed4569 100644
--- a/net/core/dev.c
+++ b/net/core/dev.c
@@ -95,6 +95,7 @@
#include <linux/notifier.h>
#include <linux/skbuff.h>
#include <linux/bpf.h>
+#include <linux/bpf_trace.h>
#include <net/net_namespace.h>
#include <net/sock.h>
#include <net/busy_poll.h>
@@ -4247,6 +4248,88 @@ static int __netif_receive_skb(struct sk_buff *skb)
    return ret;
}

+static struct static_key generic_xdp_needed __read_mostly;
+
+static int generic_xdp_install(struct net_device *dev, struct netdev_xdp *xdp)
+{
+   struct bpf_prog *new = xdp->prog;
+   int ret = 0;
+
+   switch (xdp->command) {
+   case XDP_SETUP_PROG: {
+       struct bpf_prog *old = rtnl_dereference(dev->xdp_prog);
+
+       rcu_assign_pointer(dev->xdp_prog, new);
+       if (old)
+           bpf_prog_put(old);
+
+       if (old && !new)
+           static_key_slow_dec(&generic_xdp_needed);
+       else if (new && !old)
+           static_key_slow_inc(&generic_xdp_needed);
+   }
+   }
+}

```

```

+         break;
+     }
+
+     case XDP_QUERY_PROG:
+         xdp->prog_attached = !!rcu_access_pointer(dev->xdp_prog);
+         break;
+
+     default:
+         ret = -EINVAL;
+         break;
+     }
+
+     return ret;
+}
+
+static u32 netif_receive_generic_xdp(struct sk_buff *skb,
+                                     struct bpf_prog *xdp_prog)
+{
+     struct xdp_buff xdp;
+     u32 act = XDP_DROP;
+     void *orig_data;
+     int hlen, off;
+
+     if (skb_linearize(skb))
+         goto do_drop;
+
+     /* The XDP program wants to see the packet starting at the MAC
+      * header.
+      */
+     hlen = skb_headlen(skb) + skb->mac_len;
+     xdp.data = skb->data - skb->mac_len;
+     xdp.data_end = xdp.data + hlen;
+     xdp.data_hard_start = xdp.data - skb_headroom(skb);
+     orig_data = xdp.data;
+
+     act = bpf_prog_run_xdp(xdp_prog, &xdp);
+
+     off = xdp.data - orig_data;
+     if (off)
+         __skb_push(skb, off);
+
+     switch (act) {
+     case XDP_TX:
+         __skb_push(skb, skb->mac_len);
+         /* fall through */
+     case XDP_PASS:
+         break;
+
+     default:
+         bpf_warn_invalid_xdp_action(act);
+         /* fall through */
+     case XDP_ABORTED:
+         trace_xdp_exception(skb->dev, xdp_prog, act);
+         /* fall through */
+     case XDP_DROP:
+     do_drop:
+         kfree_skb(skb);
+         break;
+     }
+
+     return act;
+}
+
+static int netif_receive_skb_internal(struct sk_buff *skb)
+{
+     int ret;
+@@ -4258,6 +4341,21 @@ static int netif_receive_skb_internal(struct sk_buff *skb)
+
+     rcu_read_lock();

```

```

+   if (static_key_false(&generic_xdp_needed)) {
+       struct bpf_prog *xdp_prog = rcu_dereference(skb->dev->xdp_prog);
+
+       if (xdp_prog) {
+           u32 act = netif_receive_generic_xdp(skb, xdp_prog);
+
+           if (act != XDP_PASS) {
+               rcu_read_unlock();
+               if (act == XDP_TX)
+                   dev_queue_xmit(skb);
+               return NET_RX_DROP;
+           }
+       }
+   }
+
+   #ifdef CONFIG_RPS
+       if (static_key_false(&rps_needed)) {
+           struct rps_dev_flow voidflow, *rflow = &voidflow;
@@ -4490,7 +4588,7 @@ static enum gro_result dev_gro_receive(struct napi_struct *napi, struct sk_buff
+           enum gro_result ret;
+           int grow;

-       if (!(skb->dev->features & NETIF_F_GRO))
+       if (netif_elide_gro(skb->dev))
+           goto normal;

+       if (skb->csum_bad)
@@ -6718,6 +6816,7 @@ EXPORT_SYMBOL(dev_change_proto_down);
+   */
+   int dev_change_xdp_fd(struct net_device *dev, int fd, u32 flags)
+   {
+       int (*xdp_op)(struct net_device *dev, struct netdev_xdp *xdp);
+       const struct net_device_ops *ops = dev->netdev_ops;
+       struct bpf_prog *prog = NULL;
+       struct netdev_xdp xdp;
@@ -6725,14 +6824,16 @@ int dev_change_xdp_fd(struct net_device *dev, int fd, u32 flags)

+       ASSERT_RTNL();

-       if (!ops->ndo_xdp)
-           return -EOPNOTSUPP;
+       xdp_op = ops->ndo_xdp;
+       if (!xdp_op || (flags & XDP_FLAGS_SKB_MODE))
+           xdp_op = generic_xdp_install;
+
+       if (fd >= 0) {
+           if (flags & XDP_FLAGS_UPDATE_IF_NOEXIST) {
+               memset(&xdp, 0, sizeof(xdp));
+               xdp.command = XDP_QUERY_PROG;

-               err = ops->ndo_xdp(dev, &xdp);
+               err = xdp_op(dev, &xdp);
+               if (err < 0)
+                   return err;
+               if (xdp.prog_attached)
@@ -6748,7 +6849,7 @@ int dev_change_xdp_fd(struct net_device *dev, int fd, u32 flags)
+               xdp.command = XDP_SETUP_PROG;
+               xdp.prog = prog;

-               err = ops->ndo_xdp(dev, &xdp);
+               err = xdp_op(dev, &xdp);
+               if (err < 0 && prog)
+                   bpf_prog_put(prog);

@@ -7789,6 +7890,7 @@ EXPORT_SYMBOL(alloc_netdev_mqs);
+   void free_netdev(struct net_device *dev)
+   {
+       struct napi_struct *p, *n;
+       struct bpf_prog *prog;

```

```

    might_sleep();
    netif_free_tx_queues(dev);
@@ -7807,6 +7909,12 @@ void free_netdev(struct net_device *dev)
    free_percpu(dev->pcpu_refcnt);
    dev->pcpu_refcnt = NULL;

+    prog = rcu_dereference(dev->xdp_prog);
+    if (prog) {
+        bpf_prog_put(prog);
+        static_key_slow_dec(&generic_xdp_needed);
+    }
+
    /* Compatibility with error handling in drivers */
    if (dev->reg_state == NETREG_UNINITIALIZED) {
        netdev_freemem(dev);
diff --git a/net/core/gro_cells.c b/net/core/gro_cells.c
index c98bbfb..814e58a 100644
--- a/net/core/gro_cells.c
+++ b/net/core/gro_cells.c
@@ -13,7 +13,7 @@ int gro_cells_receive(struct gro_cells *gcells, struct sk_buff *skb)
    struct net_device *dev = skb->dev;
    struct gro_cell *cell;

-    if (!gcells->cells || skb_cloned(skb) || !(dev->features & NETIF_F_GRO))
+    if (!gcells->cells || skb_cloned(skb) || netif_elide_gro(dev))
        return netif_rx(skb);

    cell = this_cpu_ptr(gcells->cells);
diff --git a/net/core/rtnetlink.c b/net/core/rtnetlink.c
index 58419da..958a2bf 100644
--- a/net/core/rtnetlink.c
+++ b/net/core/rtnetlink.c
@@ -896,15 +896,13 @@ static size_t rtnl_port_size(const struct net_device *dev,
    return port_self_size;
}

-static size_t rtnl_xdp_size(const struct net_device *dev)
+static size_t rtnl_xdp_size(void)
{
    size_t xdp_size = nla_total_size(0) + /* nest IFLA_XDP */
-    nla_total_size(1); /* XDP_ATTACHED */
+    nla_total_size(1) + /* XDP_ATTACHED */
+    nla_total_size(4); /* XDP_FLAGS */

    if (!dev->netdev_ops->ndo_xdp)
        return 0;
    else
        return xdp_size;
+    return xdp_size;
}

static ninline size_t if_nlmsg_size(const struct net_device *dev,
@@ -943,7 +941,7 @@ static ninline size_t if_nlmsg_size(const struct net_device *dev,
    + nla_total_size(MAX_PHYS_ITEM_ID_LEN) /* IFLA_PHYS_PORT_ID */
    + nla_total_size(MAX_PHYS_ITEM_ID_LEN) /* IFLA_PHYS_SWITCH_ID */
    + nla_total_size(IFNAMSIZ) /* IFLA_PHYS_PORT_NAME */
-    + rtnl_xdp_size(dev) /* IFLA_XDP */
+    + rtnl_xdp_size() /* IFLA_XDP */
    + nla_total_size(1); /* IFLA_PROTO_DOWN */
}
@@ -1251,23 +1249,35 @@ static int rtnl_fill_link_ifmap(struct sk_buff *skb, struct net_device *dev)

static int rtnl_xdp_fill(struct sk_buff *skb, struct net_device *dev)
{
-    struct netdev_xdp xdp_op = {};
    struct nlattr *xdp;
+    u32 xdp_flags = 0;
+    u8 val = 0;
    int err;

```

```

-     if (!dev->netdev_ops->ndo_xdp)
-         return 0;
xdp = nla_nest_start(skb, IFLA_XDP);
if (!xdp)
    return -EMSGSIZE;
-     xdp_op.command = XDP_QUERY_PROG;
-     err = dev->netdev_ops->ndo_xdp(dev, &xdp_op);
-     if (err)
-         goto err_cancel;
-     err = nla_put_u8(skb, IFLA_XDP_ATTACHED, xdp_op.prog_attached);
+     if (rcu_access_pointer(dev->xdp_prog)) {
+         xdp_flags = XDP_FLAGS_SKB_MODE;
+         val = 1;
+     } else if (dev->netdev_ops->ndo_xdp) {
+         struct netdev_xdp xdp_op = {};
+
+         xdp_op.command = XDP_QUERY_PROG;
+         err = dev->netdev_ops->ndo_xdp(dev, &xdp_op);
+         if (err)
+             goto err_cancel;
+         val = xdp_op.prog_attached;
+     }
+     err = nla_put_u8(skb, IFLA_XDP_ATTACHED, val);
+     if (err)
+         goto err_cancel;

+     if (xdp_flags) {
+         err = nla_put_u32(skb, IFLA_XDP_FLAGS, xdp_flags);
+         if (err)
+             goto err_cancel;
+     }
+     nla_nest_end(skb, xdp);
+     return 0;

```

Copyright © 2017, Eklektix, Inc.
Comments and public postings are copyrighted by their creators.
Linux is a registered trademark of Linus Torvalds